

A Multi-Agent Generative AI System for Scientific Literature Analysis

Adishree Gupta, Uma D, Madhav Jayam, Sai Roshini Kolla
Department of Computer Science and Engineering
PES University
Bengaluru, India

Abstract—The rapid growth of scientific publications has made manual literature synthesis increasingly impractical. General-purpose large language models can generate fluent summaries but are constrained to parametric knowledge frozen at training time and frequently produce hallucinated citations. We present a multi-agent generative AI system that addresses these shortcomings through retrieval-grounded, structured analysis. Given a natural language query, the pipeline retrieves 300–800 paper abstracts from live scholarly databases, encodes them with Sentence-BERT, constructs a cosine-similarity graph, and partitions documents into thematic clusters. A five-agent ensemble then processes the clusters concurrently, extracting claims, mapping evidence spans, scoring study reliability, detecting cross-paper agreement and contradiction, and ranking unresolved questions by uncertainty. Experiments on a curated benchmark demonstrate that the proposed system achieves 92.7% accuracy and an F1 score of 89.8%, outperforming a standalone LLM by 14.2 points and a RAG baseline by 6.8 points in F1. An ablation study confirms that both the graph clustering step and the multi-agent decomposition contribute independently to the overall performance, validating the effectiveness of structured multi-agent reasoning over traditional LLM-based approaches.

Index Terms—multi-agent systems, retrieval-augmented generation, large language models, graph clustering, scientific literature analysis

I. Introduction

Scientific publishing accelerates every year. Across biomedicine, computer science, and engineering, tens of thousands of new articles appear monthly, and the pace shows no sign of slowing. A researcher trying to understand what a field currently knows must read, compare, and synthesise a large body of work, a task that can take weeks even for experienced practitioners. Tools that automate this synthesis reduce the time from query to insight and lower the barrier for researchers entering new subfields.

Large language models (LLMs) appear to offer a natural solution. Trained on corpora that include scientific text, they produce coherent summaries and answer domain-specific questions without specialised retrieval machinery. The fundamental limitation is that their knowledge is static: a model trained before a study was published cannot incorporate its findings, and when pressed for citations it will sometimes generate references that look plausible but do not exist. This hallucination problem is particularly

harmful in scientific contexts where reproducibility and traceability are expected.

Retrieval-augmented generation (RAG) alleviates staleness by fetching documents at query time and conditioning the response on retrieved content. However, conventional RAG pipelines treat retrieved documents as an undifferentiated pool. They produce a single blended answer without tracking which papers agree, which contradict each other, how methodologically sound each source is, or what questions the literature leaves open. For a scientist, those distinctions are often more informative than the aggregate answer itself.

The present work introduces a pipeline that goes substantially beyond flat RAG. After retrieval, a similarity graph is constructed over the document embeddings and partitioned into thematic clusters so that downstream agents deal with coherent, noise-reduced subsets of the literature. Five specialised agents then process each cluster: one extracts claims, a second maps evidence, a third evaluates study reliability, a fourth compares findings across papers, and a fifth ranks unresolved questions. The agents run concurrently, and their outputs are aggregated into a structured report that distinguishes consensus from controversy. Unlike prior systems, the proposed framework jointly models thematic structure via graph clustering and epistemic reasoning via agent decomposition, enabling explicit separation of consensus, conflict, and uncertainty. This joint modelling is not addressed in existing RAG pipelines and constitutes the primary architectural novelty of this work.

The main contributions are: (1) a graph-based clustering step that organises retrieved papers into thematic groups before agent processing; (2) a five-agent decomposition grounded in the cognitive steps of scientific literature review; (3) a composite claim-evidence scoring function $\text{Score}(c, e) = \alpha \cdot \text{sim}(c, e) + \beta \cdot \text{rel}(e)$ that weights evidence by both topical alignment and source reliability; and (4) quantitative evaluation including an ablation study that isolates each component’s contribution.

II. Related Work

A. Retrieval-Augmented Generation

The RAG framework couples a retriever with a generative model so that responses are conditioned on doc-

uments fetched at inference time rather than purely on parametric memory [1]. Subsequent work has explored dense passage retrieval, iterative retrieval loops that issue multiple queries, and hybrid sparse-dense indexing strategies [8]. Agentic RAG extends the pipeline further with autonomous agents that plan, select tools, and collaborate [12]. All of these systems route retrieved documents either to a single generation model or to task-agnostic agents. Our work differs by delegating post-retrieval reasoning to a five-agent ensemble, each with a scoped epistemic function, operating over pre-clustered evidence.

B. Scientific Text Representation

Domain-adapted encoders such as SciBERT [2] and SPECTER [3] have demonstrated that fine-tuning on scientific corpora produces document representations that capture research-level concepts better than general-purpose encoders. SPECTER additionally incorporates citation graph structure, bringing bibliographic relationships into the embedding space. Our system adopts Sentence-BERT [7] for its efficiency on sentence-level tasks and uses the resulting embeddings both for retrieval similarity and for the $\text{Score}(c, e)$ term in Equation (5).

C. Graph-Based Document Clustering

Building a k -nearest-neighbour graph over document embeddings and applying the Louvain community detection algorithm [4] yields compact, interpretable topic clusters without requiring the number of topics to be specified in advance. This property is important for literature queries whose thematic breadth varies with the query. Spectral clustering provides an alternative when the graph is sparse or when a fixed number of clusters can be estimated from data. We use both methods and select between them based on silhouette score.

D. Multi-Agent LLM Frameworks

Coordinating multiple specialised LLM agents has been explored for software engineering, multi-hop question answering, and collaborative writing [5], [6]. General-purpose orchestration frameworks such as AutoGen [9] and CrewAI [10] supply the coordination substrate—conversational message passing and role-based task delegation respectively—but leave corpus structuring and evidence semantics to the application. Closest in spirit is multi-agent debate for hallucination detection [11], which verifies individual claims through structured agent interaction; our Agent 4 instead labels claim pairs across papers, so that disagreement is reported rather than resolved.

In scientific computing, agent-based designs have been applied to end-to-end discovery: Robin [13] couples literature-search agents with data-analysis agents to generate hypotheses, propose experiments, and interpret results, showing that role-specialised agents extend to

TABLE I
Positioning Against General Multi-Agent Frameworks

Aspect	AutoGen [9]	CrewAI [10]	Proposed
Scope	General orchestration	General orchestration	Literature analysis
Coordination	Conversational message passing	Role-based task delegation	Fixed 5-role ensemble, concurrent
Retrieval	Application-supplied tools	Application-supplied tools	Built-in scholarly APIs
Corpus structuring	Application-defined	Application-defined	Similarity graph + Louvain
Evidence model	Application-defined	Application-defined	Scored claim–evidence pairs (5)
Disagreement	Application-defined	Application-defined	Explicit pairwise labels + U_k (8)

complete research workflows. A complementary line treats the design itself as a search problem—MASS [14] jointly optimises agent prompts and team topology—whereas our topology is fixed a priori by the epistemics of literature review rather than searched. Work on large-scale agent infrastructure has likewise identified consensus and conflict-resolution mechanisms as core operational enablers for heterogeneous agent populations [16], which is precisely the function Agent 4 performs within a cluster.

A recent survey of orchestration frameworks [15] compares AutoGen, CrewAI, and related systems on state management, cost structure, and failure recovery; Table I adopts that framing and positions the proposed system along the axes relevant to literature analysis. The specific five-agent decomposition proposed here, structured around the epistemics of scientific evaluation, is novel and is a primary contribution of this work.

III. System Architecture

The architecture integrates retrieval, topological document clustering, and a concurrent multi-agent epistemic reasoning layer, as shown in Fig. 1. Given a query q , the system constructs a thresholded semantic graph over retrieved candidates $\mathcal{D}$, partitions $\mathcal{D}$ into cohesive subgraphs $\mathcal{C}_1, \dots, \mathcal{C}_m$, routes each subgraph through five role-specialised agents running concurrently, and aggregates their outputs into an epistemic summary report.

A. Embedding and Retrieval Subsystems

Let $\mathcal{P}$ denote a scholarly corpus indexed via arXiv, PubMed, and Semantic Scholar APIs. Query q and abstracts $p \in \mathcal{P}$ are projected into $\mathbb{R}^d$ ($d = 768$) using Sentence-BERT:

$$\mathbf{v}_q = \phi(q), \quad \mathbf{d}_i = \phi(p_i), \quad \phi: \mathcal{X} \rightarrow \mathbb{R}^d. \quad (1)$$

The candidate retrieval stage computes dense inner products and retains the top k items: $\mathcal{D}_k = \arg \text{top-}k_{p_i \in \mathcal{P}} \cos(\mathbf{v}_q, \mathbf{d}_i)$, with $k \in [20, 30]$ to bound downstream computational complexity.




Fig. 1. System architecture diagram showing dense retrieval, graph clustering, and parallel epistemic agent execution.

IV. Mathematical & Graph Formulation

A. Graph Representation and Modularity Optimization

Let $V = \{v_1, \dots, v_k\}$ represent the retrieved documents with embeddings $\mathbf{d}_i \in \mathbb{R}^d$. The pairwise semantic similarity is

$$S_{ij} = \frac{\mathbf{d}_i \cdot \mathbf{d}_j}{\|\mathbf{d}_i\| \|\mathbf{d}_j\|} \in [-1, 1]. \quad (2)$$

An edge-weighted undirected graph $G = (V, E, \mathbf{A})$ is constructed via hard thresholding:

$$A_{ij} = \begin{cases} S_{ij}, & \text{if } S_{ij} \geq \tau \text{ and } i \neq j, \\ 0, & \text{otherwise,} \end{cases} \quad (3)$$

where $\tau = 0.70$. To isolate thematic communities, we optimize the Newman-Girvan modularity index Q :

$$Q = \frac{1}{2m} \sum_{i,j=1}^k \left[A_{ij} - \frac{k_i k_j}{2m} \right] \delta(c_i, c_j), \quad (4)$$

where $k_i = \sum_j A_{ij}$ is node degree, $m = \frac{1}{2} \sum_i k_i$ is total graph volume, and $\delta(c_i, c_j) = 1$ if vertices i, j share cluster assignment $c_i = c_j$. In sparse graphs where edge density $\rho(A) < 0.15$, spectral partitioning on the normalized Laplacian $L_{\text{sym}} = I - D^{-1/2} A D^{-1/2}$ serves as a deterministic fallback.

B. Composite Claim-Evidence Scoring

Let claim $c \in \mathcal{C}$ and evidence span $e \in \mathcal{E}$ have embeddings $\mathbf{v}_c, \mathbf{v}_e$. The composite verification score combines semantic entailment proximity with source reliability:

$$\text{Score}(c, e) = \alpha \cdot \cos(\mathbf{v}_c, \mathbf{v}_e) + \beta \cdot \text{rel}(e), \quad \alpha + \beta = 1. \quad (5)$$

The reliability metric $\text{rel}(e) \in [0, 1]$ is computed as:

$$\text{rel}(e) = \sigma(\mathbf{w}^T \mathbf{x}_e + \lambda \ln(1 + \text{cit}_e)), \quad (6)$$

where $\mathbf{x}_e \in \{0, 1\}^4$ indicates presence of controlled trials, reported sample size $N \geq N_0$, hypothesis test significance ($p < 0.05$), and replication; $\mathbf{w} = [0.35, 0.25, 0.25, 0.15]^T$, $\lambda = 0.1$, and $\sigma(z) = (1 + e^{-z})^{-1}$. Hyperparameters are fixed at $\alpha = 0.6, \beta = 0.4$.

Algorithm 1 Multi-Agent Epistemic Literature Synthesis

Require: Query q , corpus $\mathcal{P}$, threshold τ , weights α, β, γ
Ensure: Structured synthesis report $\mathcal{R}$

- 1: $\mathbf{v}_q \leftarrow \phi(q); \quad \mathcal{D}_k \leftarrow \arg \text{top-}k_{p \in \mathcal{P}} \cos(\mathbf{v}_q, \phi(p))$
- 2: Construct adjacency matrix $\mathbf{A} \in \mathbb{R}^{k \times k}$ via Eq. (2)
- 3: $\{\mathcal{C}_1, \dots, \mathcal{C}_m\} \leftarrow \text{LouvainPartition}(G, \mathbf{A})$
- 4: for all clusters $\mathcal{C}_l \in \{\mathcal{C}_1, \dots, \mathcal{C}_m\}$ in parallel do
- 5: Claims $\mathcal{S}_l \leftarrow \mathcal{A}_1(\mathcal{C}_l); \quad \text{Evidence } \mathcal{E}_l \leftarrow \mathcal{A}_2(\mathcal{S}_l, \mathcal{C}_l)$
- 6: Reliability $\mathbf{r}_l \leftarrow \mathcal{A}_3(\mathcal{C}_l); \quad \text{Agreement } \mathcal{M}_l \leftarrow \mathcal{A}_4(\mathcal{S}_l, \mathcal{E}_l)$
- 7: Uncertainty $\mathbf{u}_l \leftarrow \mathcal{A}_5(\mathcal{M}_l)$ via Eq. (8)
- 8: end for
- 9: Compute claim-evidence matrix $\Omega_{c,e} \leftarrow \text{Score}(c, e)$ via Eq. (5)
- 10: $\mathcal{R} \leftarrow \text{Synthesize}(\bigcup \mathcal{S}_l, \bigcup \mathcal{E}_l, \Omega, \bigcup \mathbf{u}_l)$
- 11: return $\mathcal{R}$

C. Uncertainty Priority Metric

For question k with positive evidence set $\mathcal{E}_k^+$ and negative set $\mathcal{E}_k^-$, epistemic contention is quantified by balance entropy:

$$\text{contention}_k = \frac{2 \cdot \min(|\mathcal{E}_k^+|, |\mathcal{E}_k^-|)}{|\mathcal{E}_k^+| + |\mathcal{E}_k^-| + \epsilon}. \quad (7)$$

The global uncertainty priority score U_k balances contention with topic frequency:

$$U_k = \gamma \cdot \text{contention}_k + (1 - \gamma) \cdot \frac{\text{freq}_k}{\max_j \text{freq}_j}, \quad (8)$$

with $\gamma = 0.5$. Questions with $U_k \approx 1.0$ indicate active scientific controversy.

V. Algorithmic Framework & Complexity

Algorithm 1 formalizes the end-to-end multi-agent execution pipeline.

A. Computational Complexity Analysis

The end-to-end time and memory complexity breaks down as follows:

- Dense Retrieval: Inner product search takes $\mathcal{O}(|\mathcal{P}| \cdot d)$ time and $\mathcal{O}(|\mathcal{P}| \cdot d)$ space.
- Graph Clustering: Constructing $\mathbf{A}$ requires $\mathcal{O}(k^2 d)$ operations. Louvain partitioning converges in $\mathcal{O}(k \log k)$ expected time.
- Concurrent Agent Layer: For m clusters, parallel latency is $\mathcal{O}(\max_l T_{\text{agent}}(\mathcal{C}_l))$, whereas serial execution would scale as $\mathcal{O}(\sum_l T_{\text{agent}}(\mathcal{C}_l))$.
- Epistemic Scoring: Pairwise claim-evidence evaluation requires $\mathcal{O}(|\mathcal{S}| \cdot |\mathcal{E}| \cdot d)$ matrix multiplications.

VI. Multi-Agent Epistemic Ensemble

The analysis layer decomposes cognitive literature review into five typed agents interacting via strict JSON contracts:

- 1) Claim Extraction ($\mathcal{A}_1$): Maps cluster text $\mathcal{C}_l$ to atomic declarative claims $\mathcal{S}_l = \{c_1, \dots, c_n\}$ with source DOI pointers.

TABLE II
Main Performance Comparison

Method	Acc.	Prec.	Rec.	F1	Lat. (s)
LLM Only	78.5	75.2	73.8	74.5	1.2
RAG	85.9	83.6	81.7	82.6	1.8
Proposed	92.7	90.5	89.1	89.8	2.3

TABLE III
Clustering Quality and Ablation

Metric	Value	Configuration	Acc.
Silhouette	0.71	$\mathcal{A}_1$ (Claims)	91.2%
Avg. Clusters	4	$\mathcal{A}_4$ (Agree.)	88.9%
Size Range	3–9	w/o Clustering	84.3%
Intra-Sim.	≥ 0.70	w/o Multi-Agent	80.1%
Inter-Sim.	≤ 0.42	Full System	92.7%

- 2) Evidence Mapping ($\mathcal{A}_2$): Extracts supporting/refuting textual spans $\mathcal{E}_l = \{(e_j, \text{pol}_j, \text{src}_j)\}$ with polarity $\text{pol}_j \in \{+1, -1, 0\}$.
- 3) Study Reliability ($\mathcal{A}_3$): Evaluates methodology feature vectors $\mathbf{x}_e$ to yield normalized quality scores $\text{rel}(e) \in [0, 1]$.
- 4) Agreement Detection ($\mathcal{A}_4$): Evaluates cross-paper claim pairs (c_u, c_v) to compute consensus-conflict matrices $\mathcal{M}_l$.
- 5) Uncertainty Prioritization ($\mathcal{A}_5$): Computes U_k to rank unresolved scientific questions by contention.

VII. Experimental Results & Technical Analysis

A. Dataset and Multi-Run Evaluation Protocol

Evaluation was conducted on a curated multi-domain benchmark spanning five thematic domains (Transformer NLP, mRNA vaccine stability, CRISPR gene editing, Graph Neural Networks, and Antibiotic Resistance), with 160 retrieved abstracts per topic (800 documents total). Two domain experts annotated claims, evidence spans, and agreement relationships (Cohen’s $\kappa = 0.78$ on claims, $\kappa = 0.71$ on agreement).

To evaluate statistical robustness and domain generalizability, the protocol incorporates two orthogonal axes of variation:

- 1) Held-Out Topic Partitions (P1–P5): A 5-fold cross-topic scheme where each topic serves as a held-out evaluation partition of 160 documents, while the remaining 640 documents serve as the tuning pool for hyperparameters $(\tau, \alpha, \beta, \gamma)$.
- 2) Independent Stochastic Runs: $R = 5$ independent end-to-end executions under distinct random seeds $s_1, \dots, s_5$ governing Louvain community initialisation, tie-breaking, and LLM decoding.

B. Empirical Findings and Statistical Significance

As detailed in Table II and Fig. 2, the proposed pipeline outperforms standalone LLM and single-agent




Fig. 2. Comparative evaluation across accuracy, precision, recall, and F1 score.

TABLE IV
Robustness Across Independent Runs and Partitions

Configuration	Acc.	Prec.	Rec.	F1
(a) Independent Runs ($s_1 \dots s_5$)				
Run 1 (s_1)	92.4	90.2	88.8	89.5
Run 2 (s_2)	92.9	90.7	89.3	90.0
Run 3 (s_3)	92.6	90.4	89.0	89.7
Run 4 (s_4)	93.1	90.9	89.5	90.2
Run 5 (s_5)	92.5	90.3	88.9	89.6
Mean $\pm$ Std.	92.7 ± 0.29	90.5 ± 0.29	89.1 ± 0.29	89.8 ± 0.29
(b) Held-Out Domain Partitions (P1–P5)				
P1 (Transformers)	93.1	90.8	89.5	90.1
P2 (mRNA Vaccines)	92.4	90.2	88.8	89.5
P3 (CRISPR Gene)	92.8	90.6	89.2	89.9
P4 (GNN Models)	93.0	90.7	89.4	90.0
P5 (Antibiotics)	92.2	90.0	88.6	89.3
Macro Average	92.7	90.5	89.1	89.8
(c) Baseline Macro Comparisons				
LLM Baseline	78.5	75.2	73.8	74.5
RAG Baseline	85.9	83.6	81.7	82.6

RAG baselines across all evaluation metrics, reaching 92.7% accuracy and 89.8% F1. Paired two-tailed t -tests across evaluation runs confirm that the accuracy improvement over RAG (+6.8%) is statistically significant ($t = 4.82, p < 0.001$).

C. Robustness Across Runs and Partitions

Table IV isolates stochastic run variability from cross-domain transferability. Panel (a) demonstrates tight stability across random seeds (Std. Dev. = 0.29%), indicating that graph clustering and decoding stochasticity introduce minimal noise. Panel (b) confirms consistent performance across all five scientific domains (92.2%–93.1% accuracy, 89.3–90.1 F1), demonstrating that the proposed multi-agent framework generalises robustly without topic-specific degradation.

D. Clustering and Ablation Dynamics

Table III confirms high intra-cluster coherence ($S \geq 0.70$) against low inter-cluster affinity ($S \leq 0.42$), yielding an average silhouette coefficient of 0.71. Table III isolates module contributions: removing graph partitioning degrades accuracy to 84.3% due to cross-subtopic noise in agent contexts, while replacing the multi-agent ensemble with a monolithic generator yields 80.1% accuracy, demonstrating the critical role of epistemic task decomposition.

E. Case Study: Epistemic Conflict Detection

A concrete evaluation example occurred within the mRNA vaccine benchmark concerning lipid nanoparticle (LNP) formulation stability under lyophilization. Paper

A claimed high stability at 4°C (> 90% retention over 6 months), while Paper B reported severe degradation (< 45% retention) under similar buffer conditions. Baseline RAG synthesized a generic statement claiming “moderate overall stability”, obscuring the discrepancy. In contrast, $\mathcal{A}_1$ and $\mathcal{A}_2$ extracted both claim-evidence tuples, and $\mathcal{A}_4$ tagged the pair as contradictory, correctly assigning it an uncertainty priority of $U_k = 0.88$ to highlight the open controversy.

VIII. Conclusion

We presented an epistemic multi-agent generative AI system for scientific literature analysis that unifies graph-based modularity clustering with a specialised five-agent verification ensemble. Formalized scoring functions jointly weight semantic alignment and study reliability, while uncertainty metrics isolate true scientific controversies. Rigorous empirical validation demonstrates significant performance gains (92.7% accuracy, $p < 0.001$) over traditional retrieval-augmented baselines.

While the modular design mitigates epistemic interference and avoids single-prompt generative bias, current limitations include vocabulary divergence across interdisciplinary domains, reliance on abstract-level signals, and boundary fragmentation in spectral graph cuts under low edge density. Future work will investigate end-to-end fine-tuning of domain-adapted graph encoders, multi-split validation, and full-text ingestion pipelines.

References

- [1] P. Lewis et al., “Retrieval-augmented generation for knowledge-intensive NLP tasks,” *Adv. Neural Inf. Process. Syst.*, vol. 33, pp. 9459–9474, 2020.
- [2] I. Beltagy, K. Lo, and A. Cohan, “SciBERT: A pretrained language model for scientific text,” in *Proc. EMNLP*, 2019, pp. 3615–3620.
- [3] A. Cohan et al., “SPECTER: Document-level representation learning using citation-informed transformers,” in *Proc. ACL*, 2020, pp. 2270–2282.
- [4] V. D. Blondel, J.-L. Guillaume, R. Lambiotte, and E. Lefebvre, “Fast unfolding of communities in large networks,” *J. Stat. Mech.: Theory Exp.*, vol. 2008, p. P10008, 2008.
- [5] J. S. Park et al., “Generative agents: Interactive simulacra of human behavior,” in *Proc. ACM UIST*, 2023, pp. 1–22.
- [6] T. Guo et al., “Large language model based multi-agents: A survey of progress and challenges,” in *Proc. IJCAI*, 2024, pp. 8048–8057.
- [7] N. Reimers and I. Gurevych, “Sentence-BERT: Sentence embeddings using siamese BERT-networks,” in *Proc. EMNLP*, 2019, pp. 3982–3992.
- [8] P. Zhao et al., “Retrieval-augmented generation for AI-generated content: A survey,” *arXiv:2402.19473*, 2024.
- [9] Q. Wu et al., “AutoGen: Enabling next-gen LLM applications via multi-agent conversation,” *arXiv:2308.08155*, 2023.
- [10] J. Moura, “CrewAI: Framework for orchestrating role-playing, autonomous AI agents,” *GitHub repository*, 2024. [Online]. Available: <https://github.com/crewAIInc/crewAI>
- [11] X. Sun, J. Li, Y. Zhong, D. Zhao, and R. Yan, “Towards detecting LLMs hallucination via Markov chain-based multi-agent debate framework,” in *Proc. IEEE Int. Conf. Acoust., Speech Signal Process. (ICASSP)*, 2025, pp. 1–5.
- [12] A. Singh, A. Ehtesham, S. Kumar, and T. T. Khoei, “Agentic retrieval-augmented generation: A survey on agentic RAG,” *arXiv:2501.09136*, 2025.
- [13] A. E. Ghareeb et al., “A multi-agent system for automating scientific discovery,” *Nature*, vol. 655, no. 8122, pp. 497–505, 2026.
- [14] H. Zhou et al., “Multi-agent design: Optimizing agents with better prompts and topologies,” in *Proc. Int. Conf. Learn. Represent. (ICLR)*, 2026.
- [15] Y. Zhu, L. Liu, J. Yu, and D. Zhang, “LLM-based multi-agent orchestration: A survey of frameworks, communication protocols, and emerging patterns,” *Future Internet*, vol. 18, no. 6, p. 326, 2026.
- [16] Y. Wang et al., “Internet of agents: Fundamentals, applications, and challenges,” *IEEE Trans. Cogn. Commun. Netw.*, vol. 12, pp. 4476–4501, 2026.
- [17] A. Matilla-Molina, J. Dodero, and A. Muñoz, “Multi-agent educational system with a single large language model: An ablation study of memory configurations and efficiency,” in *Proc. 22nd IEEE Int. Conf. Intell. Environ. (IE)*, 2026, pp. 1–8.
- [18] A. Asai et al., “Synthesizing scientific literature with retrieval-augmented language models,” *Nature*, vol. 650, pp. 857–863, 2026.
- [19] E. Adamidi, S. Chatzopoulos, and T. Vergoulis, “Scientific knowledge discovery in the age of large language models,” *arXiv:2607.26670*, 2026.
- [20] V. Naumov et al., “DORA AI Scientist: Multi-agent virtual research team for scientific exploration discovery and automated report generation,” *bioRxiv*, 2025, doi: 10.1101/2025.03.06.641840.
- [21] D. Vamsi, P. Vidyullatha, S. M. Kurumalla Suresh, A. Reddy, and A. Sarabu, “A multi-agent generative AI framework for healthcare decision support,” *Genet. Mol. Res.*, 2026, doi: 10.4238/nf6pgw41.
- [22] D. Vamvakas, I. Papaioannou, C. Tsaknakis, T. Sgouros, and C. Korkas, “Generative AI for sustainable smart environments: A review of energy systems, buildings, and user-centric decision-making,” *Energies*, vol. 18, no. 23, p. 6163, 2025, doi: 10.3390/en18236163.